
E0:270 - Machine Learning - Neural Network Pruning Techniques

Braunstein, Cameron¹ Nair, Abhishek² Shaw, Vishal³

Abstract

This project investigated various methods of pruning a neural network. After a literature survey on various algorithms, we finalized two distinct approaches: Layered Optimal Brain Surgeon (L-OBS) and Neural Network Pruning by Significance (N2PS) for implementation. The findings were clear that L-OBS works well for larger networks compared to N2PS algorithm in terms of pruning. We would want to explore a novel algorithm for pruning based on our findings.

1. Motivation

Deep neural networks are flexible function approximators that have been very successful in a broad range of tasks. They can easily scale to millions of parameters, which results in challenges to store and load the network. Additionally, large networks can overfit to training data and can even memorize random patterns (Louizos et al., 2017). Neural network pruning seeks to remove unimportant weights from the network. This reduces the network's size in memory and can reduce overfitting.

2. Literature Review

The Optimal Brain Surgeon algorithm (OBS) takes a pre-trained network, and prunes weights based on their calculated effect on the error using a second order Taylor expansion of the error function, and adjusts the unpruned weights to compensate (Hassibi et al., 1993). Layerwise OBS (L-OBS) simplifies the OBS algorithm by only considering a weight's effect on local error. It has been demonstrated that any increase in the total error from pruning a weight is reasonably bounded above by the corresponding increase

in local error by pruning that weight (Dong et al., 2017). Thus L-OBS approximates OBS and is more computationally feasible.

The N2PS algorithm optimizes the structure of neural networks by pruning neurons from each layer. The pruning is done based on the significance of each neuron, which is calculated by getting the normalized average value of neurons for each layer, followed by comparing each neuron value with the normalized average value. The neurons which have less than significant value is pruned and the network is retrained for accuracy (Augasta & Kathirvalavakumar, 2011).

L0 regularization is a pruning technique which can be executed during the networks' initial training phase. The L0 norm of the weights is non differentiable. However by smoothing the expected L0 norm with continuous distributions with hard nonlinearity, it is possible to maintain exact zeros in the pruned weights while still allowing for efficient gradient based optimization (Louizos et al., 2017).

The Deep Compression method combines pruning with other size reduction techniques. First, a trained network is pruned by magnitude, ie, weights with an absolute value below a threshold are dropped. Next, the weights are quantized so that multiple connections share the same weight, thus only the effective weights and the indices need to be stored. Finally, apply Huffman coding to take advantage of the biased distribution of effective weights (Han et al., 2015).

3. Experiment Description

We trained a 784-300-100-10 feed forward neural network with L2 regularization to classify the MNIST dataset. There were 266200 prunable weights in total. We then tested several pruning algorithms on the network.

In our first experiment, we ran N2PS algorithm on this network.

We then compared three approaches, which could be run until they reached a compression ratio of 0. We implemented magnitude based weight pruning as a control. Against this, we ran a simplified version of the L-OBS algorithm, which recalculated the inverse Hessian only after every

¹Mathematics Department, Indian Institute of Science, Bangalore ²Department of Electronic System Engineering, Indian Institute of Science, Bangalore ³Electrical Communication Engineering, Indian Institute of Science, Bangalore. Correspondence to: Braunstein, Cameron <comeronb@iisc.ac.in>, Nair, Abhishek <abhisheknair@iisc.ac.in>, Shaw, Vishal <vishalshaw@iisc.ac.in>.


2000 pruned weights. We also ran this simplified L-OBS algorithm, with 3 iterations of retraining after every 2000 pruned weights and then a recalculation of the inverse Hessian. In the original L-OBS algorithm, the inverse Hessian is recalculated after every pruned weight. However this change was made out of computational necessity.

4. Preliminary Results

The N2PS algorithm as per literature performed the best for smaller networks like iris-dataset, breast-cancer, diabetes, ionosphere, wave, etc. These dataset are of the order of 50-10-2 network architecture. For these architecture, N2PS pruning helped improve the accuracy of the network. Results which were obtained on 748-300-100-10 network using N2PS seemed to reduce the accuracy .

depth.

One interesting algorithm we would want to explore is to try and optimize L0 regularization and implement it layer wise just like we have seen in case of OBS.



Our control, the Magnitude Based Pruning, dropped in accuracy the fastest, followed by our simplified L-OBS algorithm and then simplified L-OBS with retraining. The L-OBS with retraining graph is particularly jagged, as retraining every 2000 prunings resulted in spikes of accuracy, particularly as the network became more sparse. Interestingly, all algorithms stayed closed to their initial test accuracy until a compression ratio of approximately 0.6. This suggests that the network has redundancies which can be eliminated before training begins.

From the results, we clearly see that N2PS seems to struggle with bigger than 100 neuron network. On the other hand L-OBS seems to be a stable and viable option for neural network pruning for larger networks.

5. Future Work

We hope to combine several of our tested algorithms to see if we can achieve even better compression without loss in accuracy. We also hope to explore L0 regularization in

References

- Augasta, M. Gethsiyal and Kathirvalavakumar, T. A novel pruning algorithm for optimizing feedforward neural network of classification problems. *Neural Processing Letters*, 34:241–258, 12 2011. doi: 10.1007/s11063-011-9196-7.
- Dong, Xin, Chen, Shangyu, and Pan, Sinno Jialin. Learning to prune deep neural networks via layer-wise optimal brain surgeon. *Advances in Neural Information Processing Systems (NIPS 2017)*, 30, 11 2017. URL <https://arxiv.org/pdf/1705.07565>.
- Han, Song, Mao, Huizi, and Dally, William J. Deep compression: Compressing deep neural network with pruning, trained quantization and huffman coding. *CoRR*, 2015. URL <http://arxiv.org/abs/1510.00149>.
- Hassibi, Babak, G. Stork, David, and J. Wolff, Gregory. Optimal brain surgeon and general network pruning. pp. 293 – 299 vol.1, 02 1993.
- Louizos, C., Welling, M., and Kingma, D. P. Learning Sparse Neural Networks through L_0 Regularization. *ArXiv e-prints*, 12 2017.